

Common use of PyTorch from the basics to the advanced

Daesoo Lee

NTNU



Norwegian University of
Science and Technology

05.04.2022
NTNU

Few Remarks

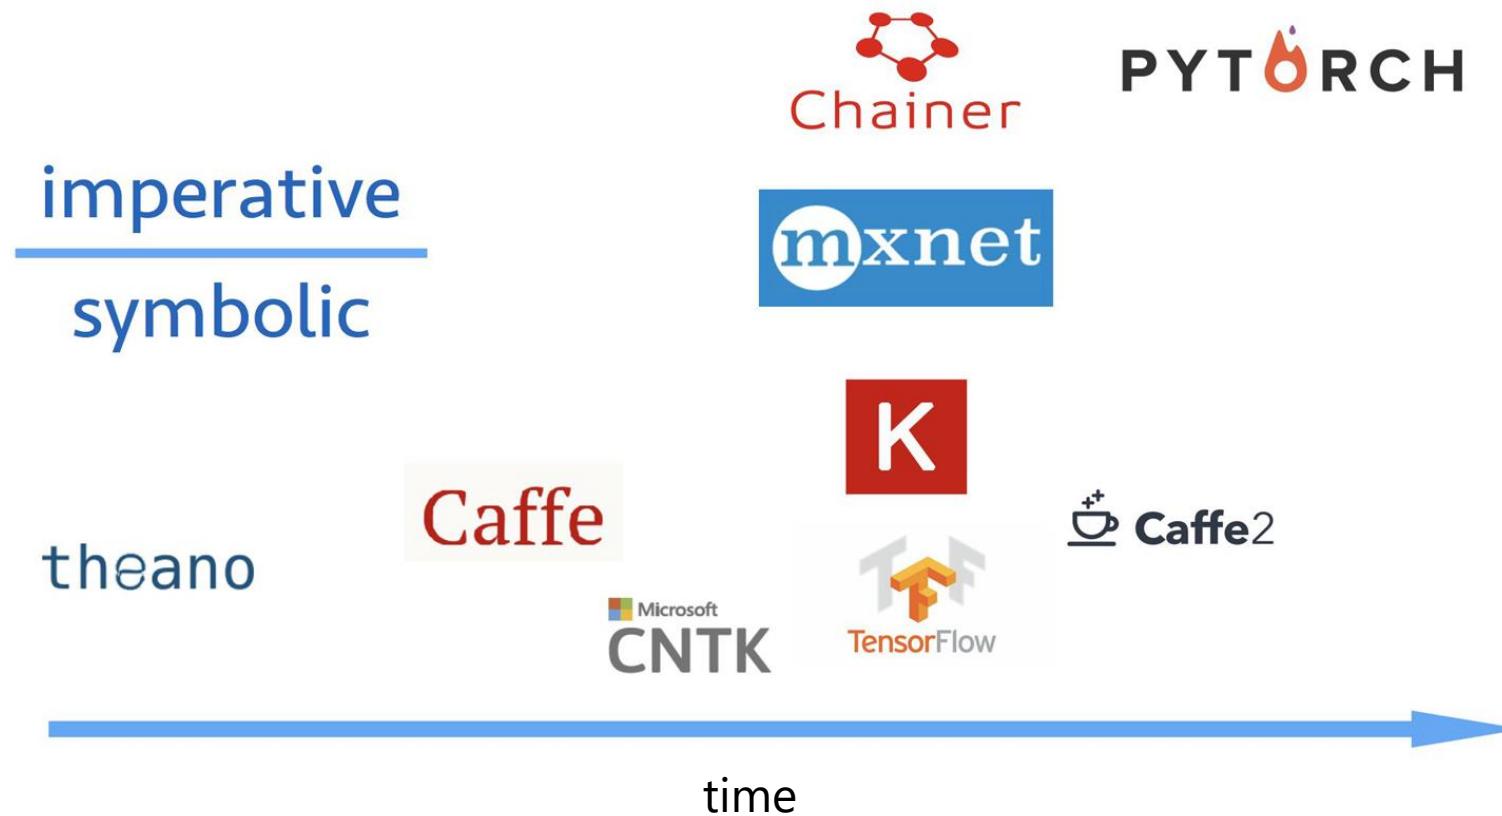
This seminar assumes that

- You have a basic knowledge about neural network stuff (*e.g.*, basic neural network, gradient descent).
- You've heard of PyTorch and knows what it's for.

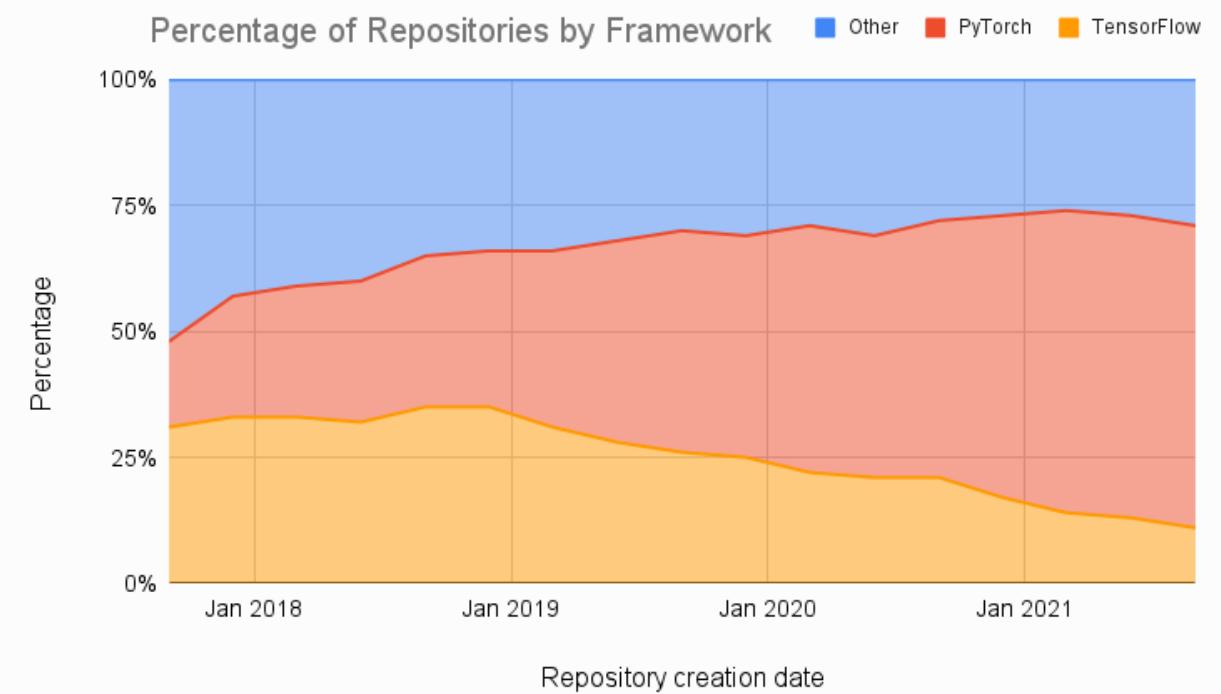
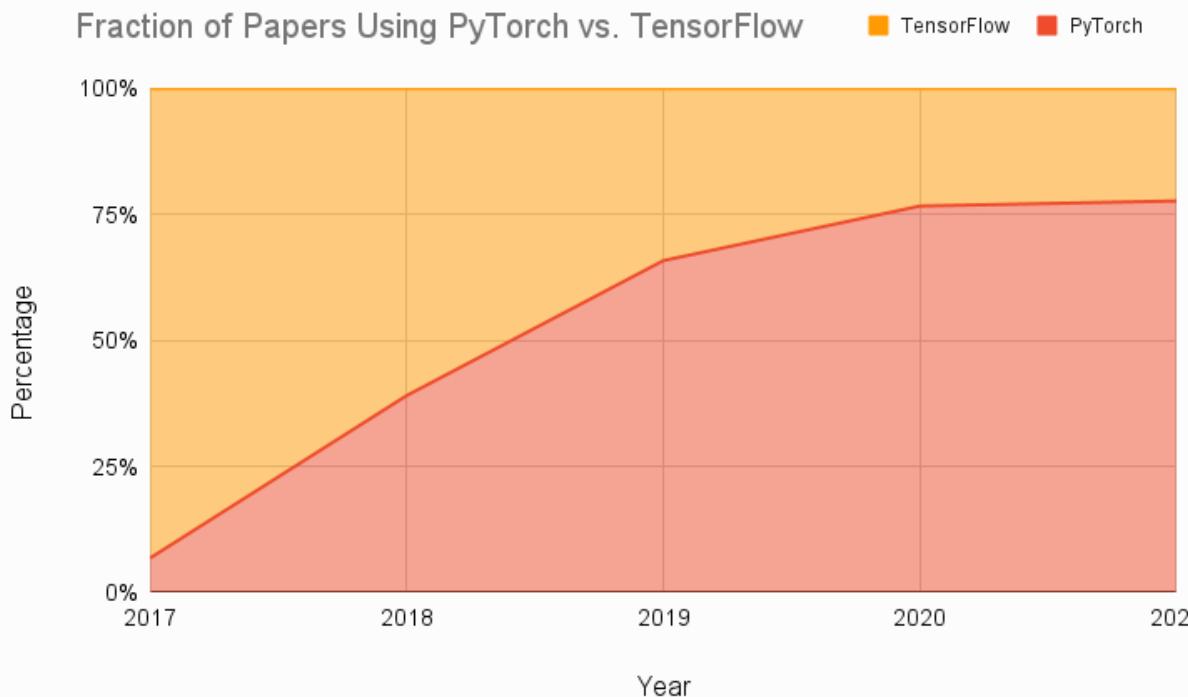
Why PyTorch?

Why PyTorch?

Popular Deep Learning Frameworks

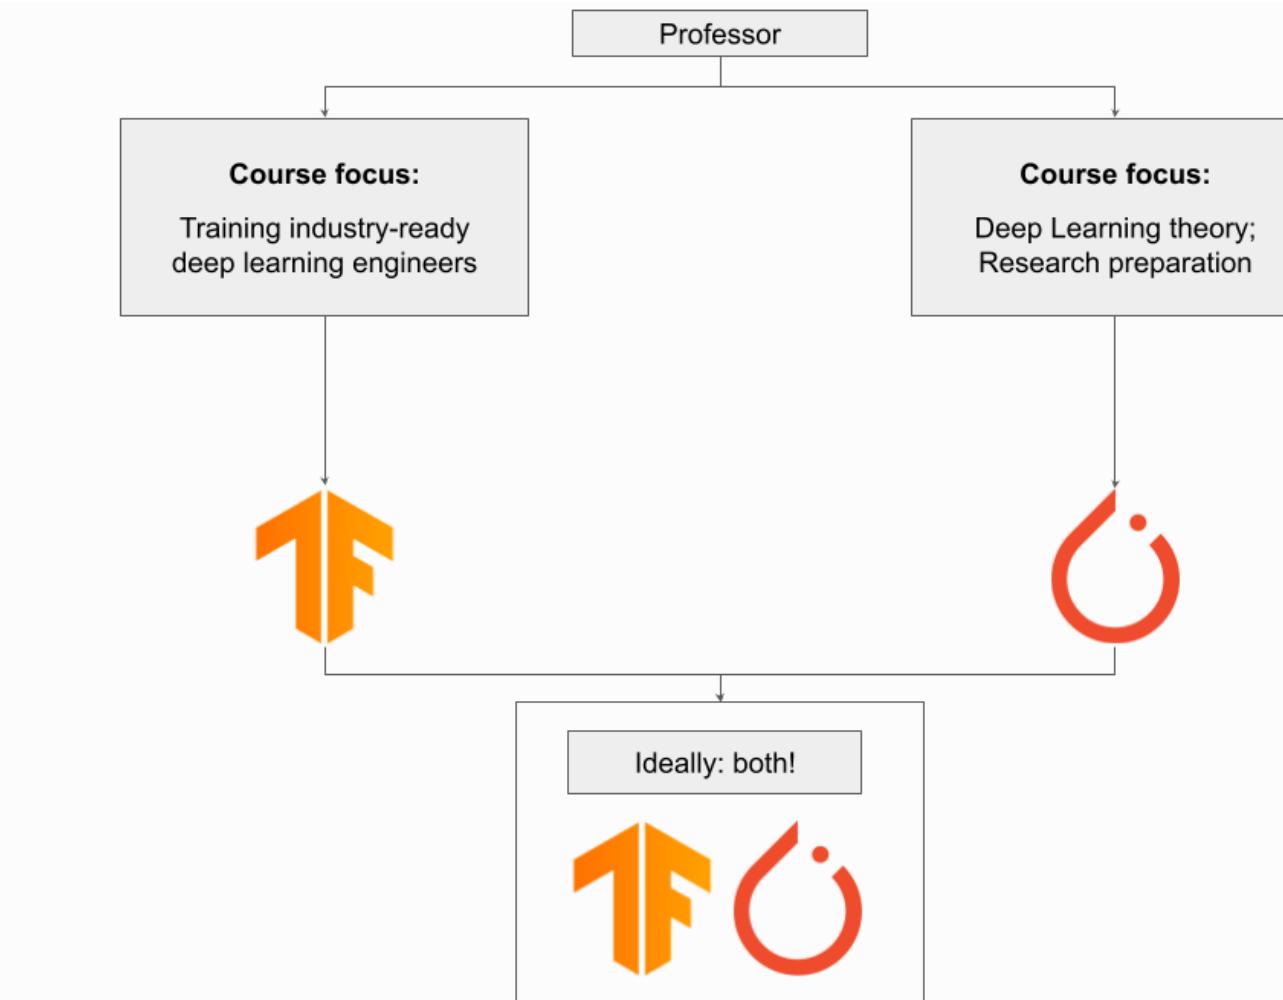


Why PyTorch?



Why PyTorch?

What if You're a Professor? (if there's any here today)

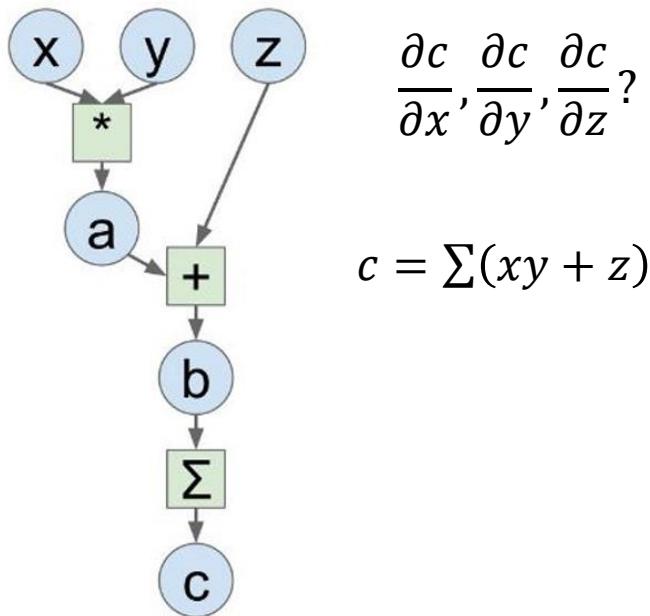


Why PyTorch?

	Keras 	TensorFlow 	PyTorch 
Level of API	high-level API ¹	Both high & low level APIs	Lower-level API ²
Speed	Slow	High	High
Architecture	Simple, more readable and concise	Not very easy to use	Complex ³
Debugging	No need to debug	Difficult to debugging	Good debugging capabilities
Dataset Compatibility	Slow & Small	Fast speed & large	Fast speed & large datasets
Popularity Rank	1	2	3
Uniqueness	Multiple back-end support	Object Detection Functionality	Flexibility & Short Training Duration
Created By	Not a library on its own	Created by Google	Created by Facebook ⁴
Ease of use	User-friendly	Incomprehensive API	Integrated with Python language
Computational graphs used	Static graphs	Static graphs	Dynamic computation graphs ⁵

Why PyTorch?

Computation Graph



Tensorflow

```
import numpy as np
np.random.seed(0)
import tensorflow as tf

N, D = 3, 4

with tf.device('/gpu:0'):
    x = tf.placeholder(tf.float32)
    y = tf.placeholder(tf.float32)
    z = tf.placeholder(tf.float32)

    a = x * y
    b = a + z
    c = tf.reduce_sum(b)

grad_x, grad_y, grad_z = tf.gradients(c, [x, y, z])

with tf.Session() as sess:
    values = {
        x: np.random.randn(N, D),
        y: np.random.randn(N, D),
        z: np.random.randn(N, D),
    }
    out = sess.run([c, grad_x, grad_y, grad_z],
                  feed_dict=values)
    c_val, grad_x_val, grad_y_val, grad_z_val = out
```

PyTorch

```
import torch

N, D = 3, 4

x = torch.rand((N, D), requires_grad=True)
y = torch.rand((N, D), requires_grad=True)
z = torch.rand((N, D), requires_grad=True)

a = x * y
b = a + z
c = torch.sum(b)

c.backward()
```

Quick Computation Speed Experiment for Matrix Multiplication

Quick Computation Speed Experiment for Matrix Multiplication



NumPy

```
import numpy as np
import torch

# task: compute matrix multiplication C = AB
d = 3000

# using numpy
A = np.random.rand(d, d).astype(np.float32)
B = np.random.rand(d, d).astype(np.float32)
C = A.dot(B) ←
```

350 ms

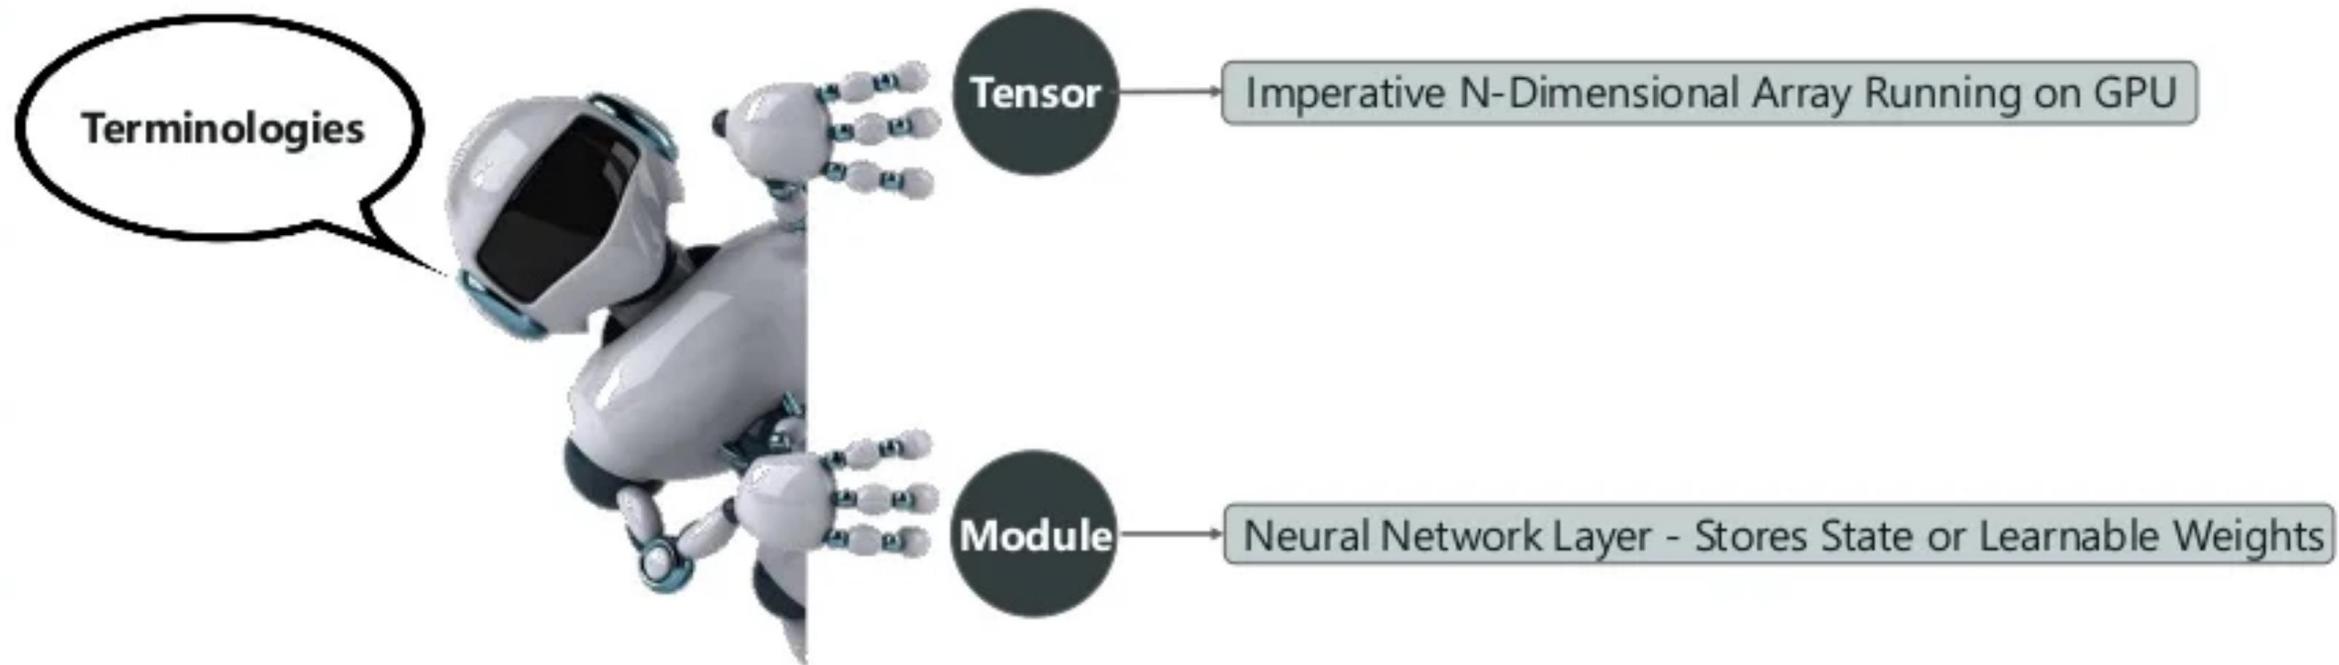
PYTORCH

```
# using torch with gpu
A = torch.rand(d, d).cuda()
B = torch.rand(d, d).cuda()
C = torch.mm(A, B) ←
```

0.1 ms

Major PyTorch Concepts

Major PyTorch Concepts



Tensors

Tensors are similar to NumPy's ndarrays, with the addition being that Tensors can also be used on a GPU to accelerate computing.

Common operations for creation and manipulation of these Tensors are similar to those for ndarrays in NumPy. (rand, ones, zeros, indexing, slicing, reshape, transpose, cross product, matrix product, element wise multiplication)

$$w_{t+1} = w_t - \eta \nabla w_t$$

Major PyTorch Concepts

Tensor Example: single variable

Let's see the following example: We're trying to compute a derivative of y w.r.t x :

$$y(x) = x^2$$

If $x = 2$, then

$$y(2) = 2^2 = 4$$

$$\frac{dy(x)}{dx} = 2x$$

$$\frac{dy(2)}{dx} = 2(2) = 4$$

Let's see how we can compute them with pytorch:

```
import torch
```

```
x = torch.tensor(2., requires_grad=True)
```

allows to track gradients

```
y = x**2
```

= making it a trainable parameter

```
y.backward() Backpropagation = automatic differentiation
```

```
x.grad
```

```
tensor(4.)
```

$$y(x) = x^3$$

$$\frac{dy(x)}{dx} = 3x^2$$

$$\frac{dy(2)}{dx} = 12$$

```
import torch
```

```
x = torch.tensor(2., requires_grad=True)
```

```
y = x**3
```

```
y.backward()
```

```
x.grad
```

```
tensor(12.)
```

Major PyTorch Concepts

Tensor Example: matrices

Attributes of a tensor 't':

- `t = torch.randn(1)`

`requires_grad`- making a trainable parameter

- By default False
- Turn on:
 - `t.requires_grad_()` or
 - `t = torch.randn(1, requires_grad=True)`
- Accessing tensor value:
 - `t.data`
- Accessing tensor gradient
 - `t.grad`

```
1 import torch
2
3 N, D = 3, 4
4
5 x = torch.rand((N, D), requires_grad=True)
6 y = torch.rand((N, D), requires_grad=True)
7 z = torch.rand((N, D), requires_grad=True)
8
9 a = x * y
10 b = a + z
11 c = torch.sum(b)
12
13 c.backward()
14
15
16 print(x.data)
17 print(x.grad)
```

```
tensor([[0.4118, 0.2576, 0.3470, 0.0240],
       [0.7797, 0.1519, 0.7513, 0.7269],
       [0.8572, 0.1165, 0.8596, 0.2636]])
tensor([[0.6855, 0.9696, 0.4295, 0.4961],
       [0.3849, 0.0825, 0.7400, 0.0036],
       [0.8104, 0.8741, 0.9729, 0.3821]])
```

Loading Data, Devices and CUDA

Numpy arrays to PyTorch tensors

- `torch.from_numpy(x_train)`
- Returns a cpu tensor!

PyTorch tensor to numpy

- `t.numpy()`

Using GPU acceleration

- `t.to()`
- Sends to whatever device (cuda or cpu)

Fallback to cpu if gpu is unavailable:

- `torch.cuda.is_available()`

Check cpu/gpu tensor OR numpyarray ?

- `type(t) or t.type() returns`
 - `numpy.ndarray`
 - `torch.Tensor`
 - CPU - `torch.cpu.FloatTensor`
 - GPU - `torch.cuda.FloatTensor`

```
x = np.random.rand(2, 2)
x
array([[0.18229139, 0.46024418],
       [0.32090798, 0.9557614 ]])
```

```
torch.cuda.is_available() NVIDIA GeForce RTX 3060
```

True

```
device = torch.device(0)
device
device(type='cuda', index=0)
```

```
x = torch.from_numpy(x)
x
tensor([[0.1823, 0.4602],
       [0.3209, 0.9558]], dtype=torch.float64)
```

```
x.to(device)
tensor([[0.1823, 0.4602],
       [0.3209, 0.9558]], device='cuda:0', dtype=torch.float64)
```

Major PyTorch Concepts

Model

In PyTorch, a model is represented by a regular Python class that inherits from the Module class.

- Two components

- `__init__(self)` : it defines the parts that make up the model- in our case, two parameters, `a` and `b`
- `forward(self, x)` : it performs the actual computation, that is, it outputs a prediction, given the input `x`

```
import torch.nn as nn

class ManualLinearRegression(nn.Module):
    def __init__(self):
        super().__init__()
        self.a = nn.Parameter(torch.randn(1, requires_grad=True))
        self.b = nn.Parameter(torch.randn(1, requires_grad=True))

    def forward(self, x):
        return self.a + self.b * x

reg = ManualLinearRegression()
reg

ManualLinearRegression()

for name, param in reg.named_parameters():
    print(name, param)

a Parameter containing:
tensor([0.0038], requires_grad=True)
b Parameter containing:
tensor([-0.4261], requires_grad=True)

x = torch.Tensor([1])
out = reg(x)
print(out)

tensor([-0.4223], grad_fn=<AddBackward0>)
```

Major PyTorch Concepts

```
import torch.nn as nn

class ManualLinearRegression(nn.Module):
    def __init__(self):
        super().__init__()
        self.a = nn.Parameter(torch.randn(1, requires_grad=True))
        self.b = nn.Parameter(torch.randn(1, requires_grad=True))

    def forward(self, x):
        return self.a + self.b * x

reg = ManualLinearRegression()
reg

ManualLinearRegression()

for name, param in reg.named_parameters():
    print(name, param)

a Parameter containing:
tensor([ 0.0038], requires_grad=True)
b Parameter containing:
tensor([-0.4261], requires_grad=True)

x = torch.Tensor([1])
out = reg(x)
print(out)

tensor([-0.4223], grad_fn=<AddBackward0>)
```

```
linear = nn.Linear(1, 1) # (in_size, out_size)

x = torch.Tensor([1])
out = linear(x)

print(out)

tensor([-1.5249], grad_fn=<AddBackward0>)
```

```
class LinearRegression(nn.Module):
    def __init__(self):
        super().__init__()
        self.linear1 = nn.Linear(1, 1)
        self.linear2 = nn.Linear(1, 1)

    def forward(self, x):
        return self.linear2(self.linear1(x))

reg = LinearRegression()
reg

LinearRegression(
  (linear1): Linear(in_features=1, out_features=1, bias=True)
  (linear2): Linear(in_features=1, out_features=1, bias=True)
)

x = torch.Tensor([1])
out = reg(x)
print(out)

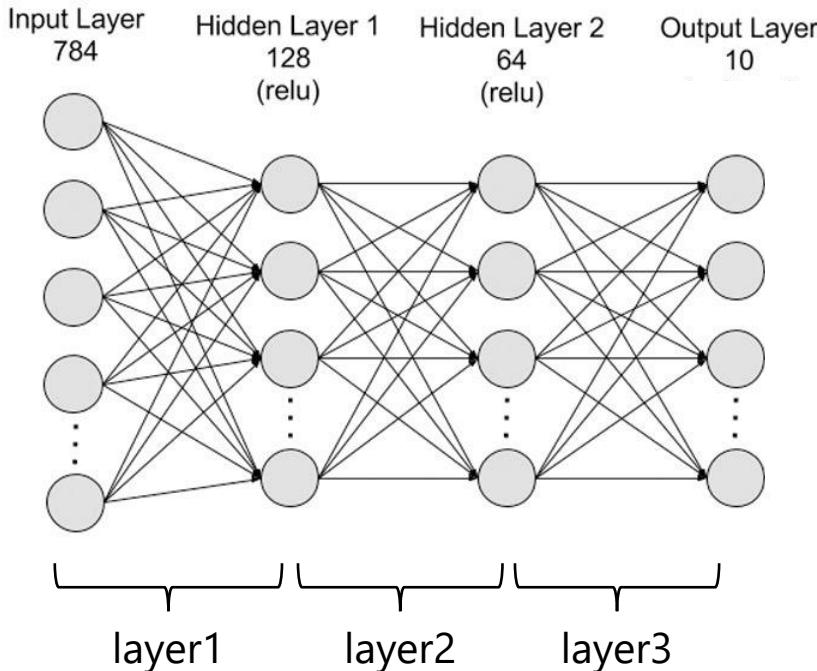
tensor([-0.3079], grad_fn=<AddBackward0>)
```

Essentially, “Module/Model” is a sort of wrapper.



Creating Simple Neural Network (Multi-Layer Perceptron)

Creating Simple Neural Network (Multi-Layer Perceptron)



```
from torch import relu

class SimpleNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.layer1 = nn.Linear(784, 128)
        self.layer2 = nn.Linear(128, 64)
        self.layer3 = nn.Linear(64, 10)

    def forward(self, x):
        out = relu(self.layer1(x))
        out = relu(self.layer2(out))
        out = self.layer3(out)
        return out

model = SimpleNN()
model
```

SimpleNN(
 (layer1): Linear(in_features=784, out_features=128, bias=True)
 (layer2): Linear(in_features=128, out_features=64, bias=True)
 (layer3): Linear(in_features=64, out_features=10, bias=True))

```
batch_size = 32
n_predictors = 784
x = torch.rand(batch_size, n_predictors)

out = model(x)
print(out.shape)
torch.Size([32, 10])
```

```
import torch.nn as nn

class SimpleNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.model = nn.Sequential(nn.Linear(784, 128), nn.ReLU(),
                                  nn.Linear(128, 64), nn.ReLU(),
                                  nn.Linear(64, 10))

    def forward(self, x):
        return self.model(x)

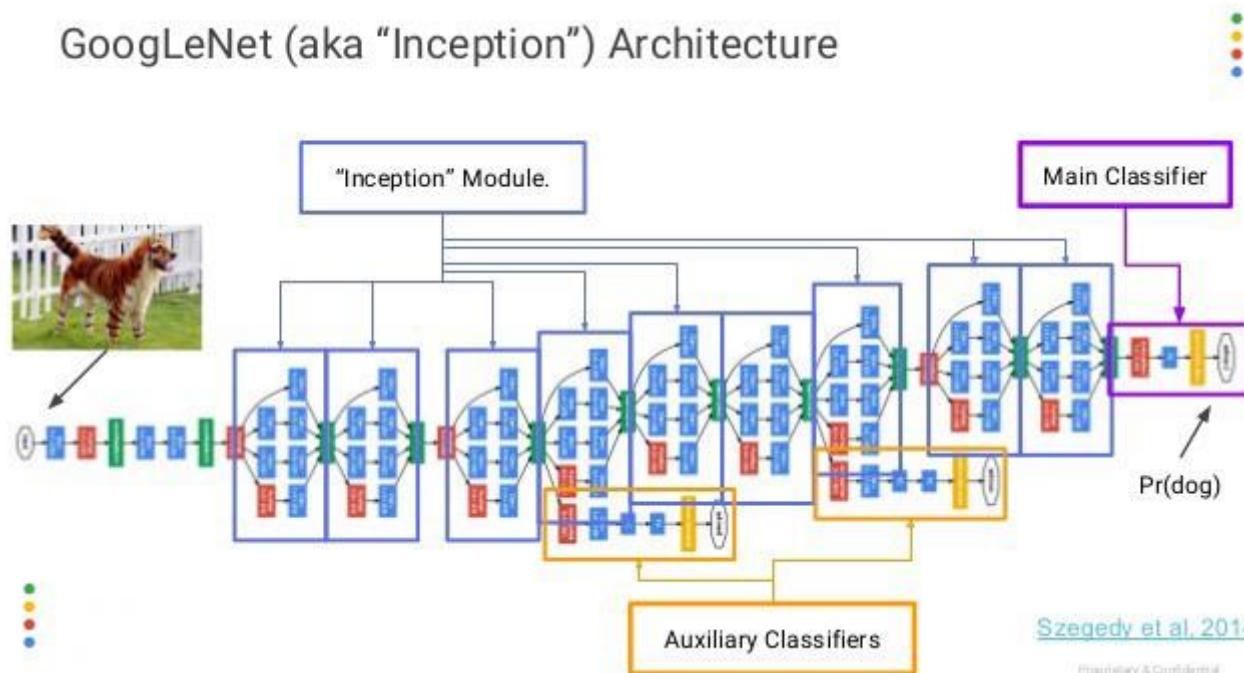
model = nn.Sequential(nn.Linear(784, 128), nn.ReLU(),
                      nn.Linear(128, 64), nn.ReLU(),
                      nn.Linear(64, 10))

out = model(x)
print(out.shape)
torch.Size([32, 10])
```

Creating Simple Neural Network (Multi-Layer Perceptron)

You'd need **nn.Module** if you're building some complex architecture

For instance,



Data Pipeline

Data Pipeline

Toy Dataset: Iris

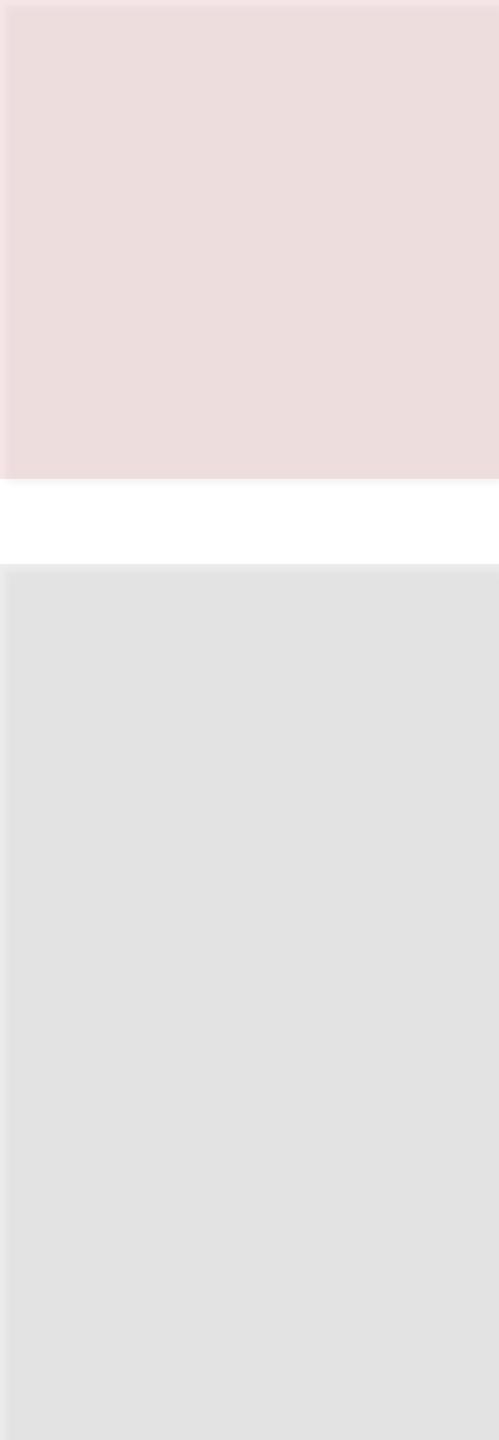
- The iris dataset is a classic and very easy multi-class classification dataset.

Classes	3
Samples per class	50
Samples total	150
Dimensionality	4
Features	real, positive

```
from sklearn.datasets import load_iris
```

```
data = load_iris()
```

```
data['target']
```

Optimizer (for Gradient Descent)

Optimizer (for Gradient Descent)

```
model = nn.Sequential(nn.Linear(784, 128), nn.ReLU(),
                      nn.Linear(128, 64), nn.ReLU(),
                      nn.Linear(64, 10))
```

```
from torch import optim
optimizer = optim.SGD(model.parameters(), lr= 0.01) # lr: Learning rate
optimizer
Register model's parameters in the optimizer.
```

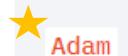
```
SGD (
Parameter Group 0
  dampening: 0
  lr: 0.01
  momentum: 0
  nesterov: False
  weight_decay: 0
)
```

SGD: Stochastic Gradient Descent

Adadelta

Adagrad

LBFGS



Adam

NAdam



AdamW

RAdam

SparseAdam

RMSprop

Adamax

Rprop

ASGD

SGD

Loss Functions

Loss Functions

- Regression: Mean Squared Loss (MSE)
- Classification: Cross Entropy Loss
- etc.

MSELOSS ⚡

```
CLASS torch.nn.MSELoss(size_average=None, reduce=None, reduction='mean') [SOURCE]
```

CROSSENTROPYLOSS

```
CLASS torch.nn.CrossEntropyLoss(weight=None, size_average=None, ignore_index=-100,  
reduce=None, reduction='mean', label_smoothing=0.0) [SOURCE]
```

Entire Training Pipeline

Entire Training Pipeline

```
from torch.utils.data import Dataset
from torch.utils.data import DataLoader
from sklearn.preprocessing import StandardScaler

class ToyDataset(Dataset):
    def __init__(self):
        data = load_iris()

        self.X = data['data']
        self.Y = data['target']

        # scale * scaling added
        scaler = StandardScaler()
        self.X = scaler.fit_transform(self.X)

    def __getitem__(self, idx):
        """
        set such that an individual sample is returned.
        """
        x = self.X[idx, :]
        y = self.Y[idx]
        return x, y

    def __len__(self):
        """
        an entire number of samples should be returned.
        """
        return self.Y.size
```

Classes	3
Samples per class	50
Samples total	150
Dimensionality	4
Features	real, positive

```
in_size = 4 # num_predictors
h_size = 32 # hidden layer size
out_size = 3 # num_classes

model = nn.Sequential(nn.Linear(in_size, h_size), nn.ReLU(),
                      nn.Linear(h_size, h_size), nn.ReLU(),
                      nn.Linear(h_size, out_size))
```

```
from torch import optim
optimizer = optim.Adam(model.parameters(),
                       lr=0.001) # lr: learning rate
optimizer
```

```
loss_func = nn.CrossEntropyLoss()
```

```
# Training
n_epochs = 20

step = 1
for epoch in range(n_epochs):

    for batch in data_loader:
        x, y = batch

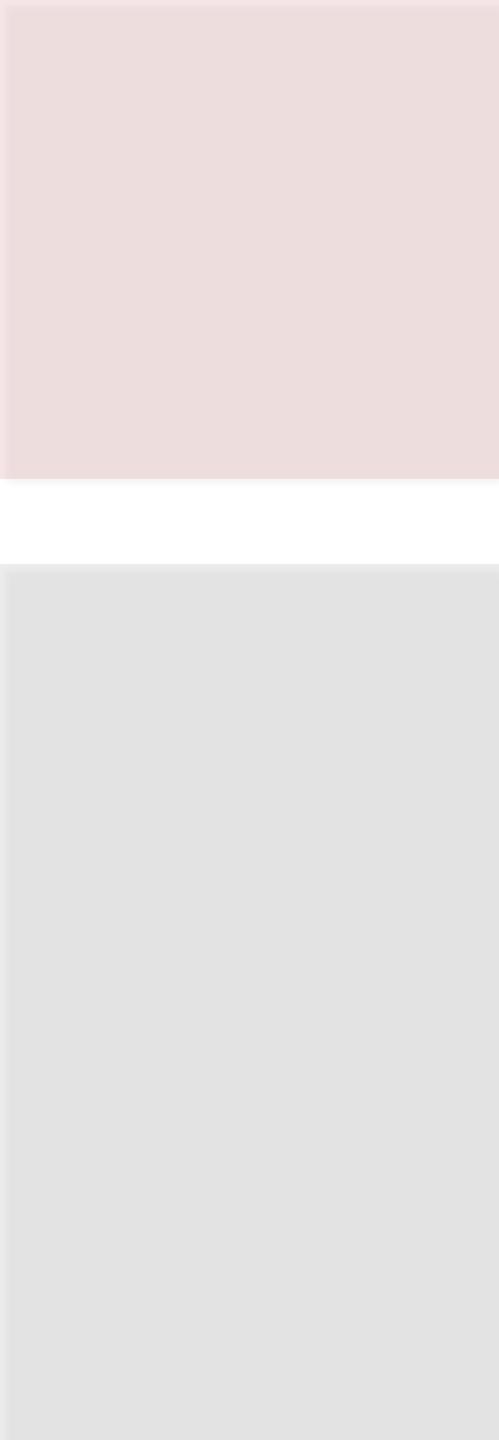
        yhat = model(x.float())
        loss = loss_func(yhat, y.long())

        optimizer.zero_grad() # reset the gradients (to zero)
        loss.backward() # compute gradients
        optimizer.step() # update model's weight

        print(f'step: {step} | loss: {loss.item()}')
        step += 1
```

Training
pipeline
template

step: 1	loss: 1.1227813959121704
step: 2	loss: 1.1198639869689941
step: 3	loss: 1.1341817378997803
step: 4	loss: 1.107405424118042
step: 5	loss: 1.1096519231796265
step: 96	loss: 0.4173617660999298
step: 97	loss: 0.4993354082107544
step: 98	loss: 0.37763741612434387
step: 99	loss: 0.4581018388271332
step: 100	loss: 0.31296491622924805



How to deal with
a large dataset?

How to deal with a large dataset?

```
from torch.utils.data import Dataset
from torch.utils.data import DataLoader
from sklearn.preprocessing import StandardScaler

class ToyDataset(Dataset):
    def __init__(self):
        data = load_iris()

        self.X = data['data']
        self.Y = data['target']

        # scale
        scaler = StandardScaler()
        self.X = scaler.fit_transform(self.X)

    def __getitem__(self, idx):
        """
        set such that an individual sample is returned.
        """
        x = self.X[idx, :]
        y = self.Y[idx]
        return x, y

    def __len__(self):
        """
        an entire number of samples should be returned.
        """
        return self.Y.size
```

```
class Dataset(Dataset):
    def __init__(self, csv_file, data_dir, transform=None):
        self.transform = transform
        self.data_dir = data_dir
        data_dircsv_file = os.path.join(self.data_dir, csv_file)
        self.data_name = pd.read_csv(data_dircsv_file)

        self.len = self.data_name.shape[0]

    def __len__(self):
        return self.len

    def __getitem__(self, idx):
        img_name = os.path.join(self.data_dir, self.data_name.iloc[idx, 1])
        image = Image.open(img_name)
        y = self.data_name.iloc[idx, 0]

        if self.transform:
            image = self.transform(image)
        return image, y
```

To load the following csv file :

	category	image
0	Ankle boot	img/fashion0.png
1	T-shirt	img/fashion1.png
2	T-shirt	img/fashion2.png
3	Dress	img/fashion3.png
4	T-shirt	img/fashion4.png

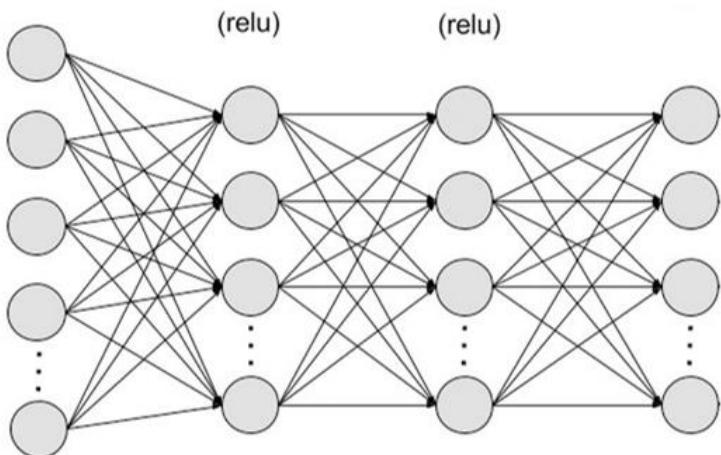
This allows data to be fetched by one by one without loading all the images into the memory.



CNN, RNN

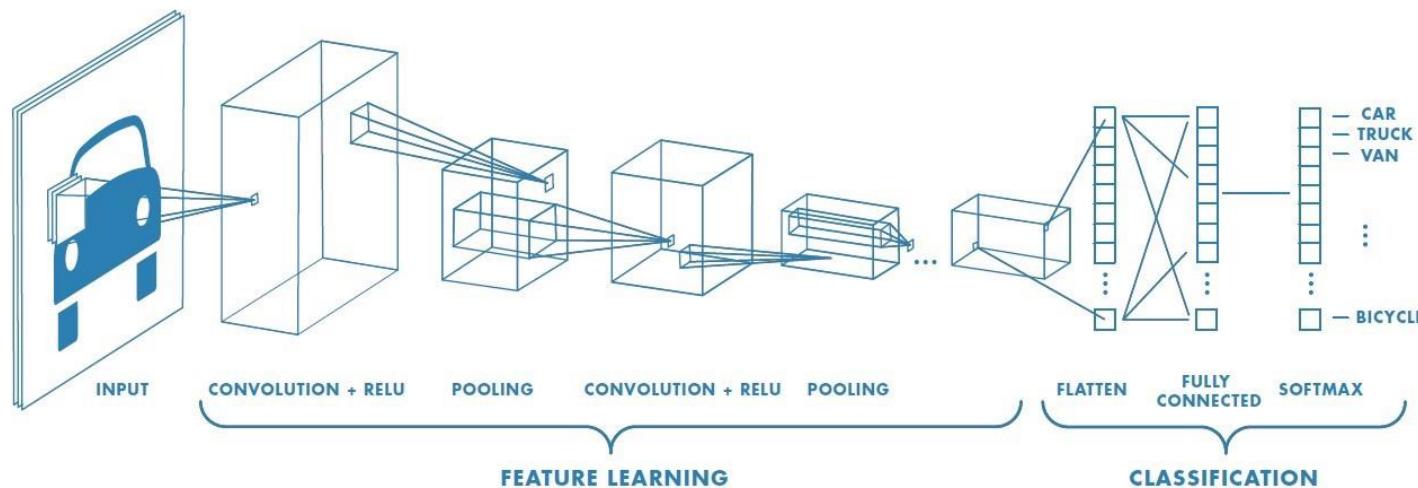
CNN, RNN

Input Layer Hidden Layer 1 Hidden Layer 2 Output Layer



```
in_size = 4 # num_predictors
h_size = 32 # hidden layer size
out_size = 3 # num_classes

model = nn.Sequential(nn.Linear(in_size, h_size), nn.ReLU(),
                      nn.Linear(h_size, h_size), nn.ReLU(),
                      nn.Linear(h_size, out_size))
```



```
in_channels = 3
model = nn.Sequential(nn.Conv2D(in_channels, out_channels=32, kernel_size=3, stride=1),
                      nn.ReLU(),
                      nn.MaxPool2d(kernel_size=3, stride=2),

                      nn.Conv2D(in_channels=32, out_channels=64, kernel_size=3, stride=1),
                      nn.ReLU(),
                      nn.MaxPool2d(kernel_size=3, stride=2),

                      nn.Conv2D(in_channels=64, out_channels=128, kernel_size=3, stride=1),
                      nn.ReLU(),
                      nn.MaxPool2d(kernel_size=3, stride=2),

                      nn.Flatten(start_dim=1),

                      nn.Linear(...),
                      nn.ReLU(),
                      nn.Linear(...))
```

The `torchvision.models` subpackage contains definitions of models for addressing different tasks, including: image classification, pixelwise semantic segmentation, object detection, instance segmentation, person keypoint detection, video classification, and optical flow.

Pretrained Models

• NOTE

Backward compatibility is guaranteed for loading a serialized `state_dict` to the model created using old PyTorch version. On the contrary, loading entire saved models or serialized `ScriptModules` (serialized using older versions of PyTorch) may not preserve the historic behaviour. Refer to the following [documentation](#)

Classification

The models subpackage contains definitions for the following model architectures for image classification:

- [AlexNet](#)
- [VGG](#)
- [ResNet](#)
- [SqueezeNet](#)
- [DenseNet](#)
- [Inception v3](#)
- [GoogLeNet](#)
- [ShuffleNet v2](#)
- [MobileNetV2](#)
- [MobileNetV3](#)
- [ResNeXt](#)
- [Wide ResNet](#)
- [MNASNet](#)
- [EfficientNet](#)
- [RegNet](#)
- [VisionTransformer](#)
- [ConvNeXt](#)



Trained on lots of images

```
import torchvision.models as models
resnet18 = models.resnet18(pretrained=True)
alexnet = models.alexnet(pretrained=True)
squeezenet = models.squeezezenet1_0(pretrained=True)
vgg16 = models.vgg16(pretrained=True)
densenet = models.densenet161(pretrained=True)
inception = models.inception_v3(pretrained=True)
googlenet = models.googlenet(pretrained=True)
shufflenet = models.shufflenet_v2_x1_0(pretrained=True)
mobilenet_v2 = models.mobilenet_v2(pretrained=True)
mobilenet_v3_large = models.mobilenet_v3_large(pretrained=True)
mobilenet_v3_small = models.mobilenet_v3_small(pretrained=True)
resnext50_32x4d = models.resnext50_32x4d(pretrained=True)
wide_resnet50_2 = models.wide_resnet50_2(pretrained=True)
mnasnet = models.mnasnet1_0(pretrained=True)
efficientnet_b0 = models.efficientnet_b0(pretrained=True)
efficientnet_b1 = models.efficientnet_b1(pretrained=True)
efficientnet_b2 = models.efficientnet_b2(pretrained=True)
efficientnet_b3 = models.efficientnet_b3(pretrained=True)
efficientnet_b4 = models.efficientnet_b4(pretrained=True)
efficientnet_b5 = models.efficientnet_b5(pretrained=True)
efficientnet_b6 = models.efficientnet_b6(pretrained=True)
efficientnet_b7 = models.efficientnet_b7(pretrained=True)
regnet_y_400mf = models.regnet_y_400mf(pretrained=True)
regnet_y_800mf = models.regnet_y_800mf(pretrained=True)
regnet_y_1_6gf = models.regnet_y_1_6gf(pretrained=True)
regnet_y_3_2gf = models.regnet_y_3_2gf(pretrained=True)
regnet_y_8gf = models.regnet_y_8gf(pretrained=True)
regnet_y_16gf = models.regnet_y_16gf(pretrained=True)
regnet_y_32gf = models.regnet_y_32gf(pretrained=True)
regnet_x_400mf = models.regnet_x_400mf(pretrained=True)
regnet_x_800mf = models.regnet_x_800mf(pretrained=True)
regnet_x_1_6gf = models.regnet_x_1_6gf(pretrained=True)
regnet_x_3_2gf = models.regnet_x_3_2gf(pretrained=True)
regnet_x_8gf = models.regnet_x_8gf(pretrained=True)
regnet_x_16gf = models.regnet_x_16gf(pretrained=True)
regnet_x_32gf = models.regnet_x_32gf(pretrained=True)
vit_b_16 = models.vit_b_16(pretrained=True)
vit_b_32 = models.vit_b_32(pretrained=True)
vit_l_16 = models.vit_l_16(pretrained=True)
vit_l_32 = models.vit_l_32(pretrained=True)
convnext_tiny = models.convnext_tiny(pretrained=True)
convnext_small = models.convnext_small(pretrained=True)
convnext_base = models.convnext_base(pretrained=True)
convnext_large = models.convnext_large(pretrained=True)
```

RNN, LSTM, GRU

RNN

CLASS `torch.nn.RNN(*args, **kwargs)` [\[SOURCE\]](#)

GRU

CLASS `torch.nn.GRU(*args, **kwargs)` [\[SOURCE\]](#)

LSTM

CLASS `torch.nn.LSTM(*args, **kwargs)` [\[SOURCE\]](#)

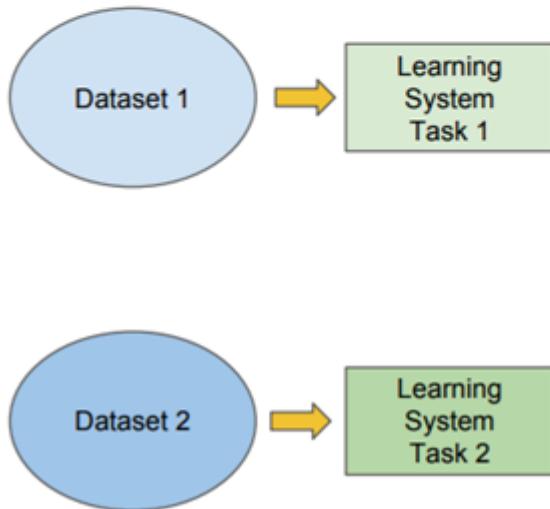
```
model = nn.Sequential(nn.LSTM(...),  
                      ...)
```

Transfer Learning

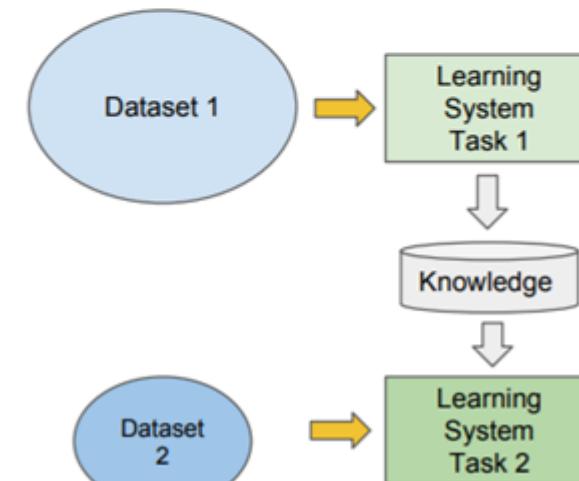
Transfer Learning

Traditional ML vs Transfer Learning

- Isolated, single task learning:
 - Knowledge is not retained or accumulated. Learning is performed w.o. considering past learned knowledge in other tasks



- Learning of a new tasks relies on the previous learned tasks:
 - Learning process can be faster, more accurate and/or need less training data



Transfer Learning

MODELS AND DATASETS

The `torchvision.models` subpackage includes: image classification, semantic segmentation, person keypoint detection.

• NOTE

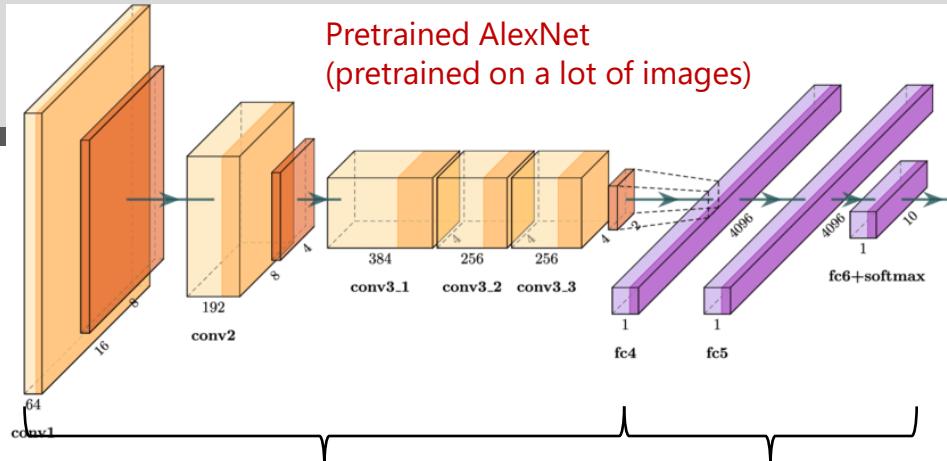
Backward compatibility is guaranteed when using old PyTorch version. Consider using `ScriptModules` (serialized under `script`) if you want to change behaviour. Refer to the following documentation for more information.

Classification

The `models` subpackage contains models for classification:

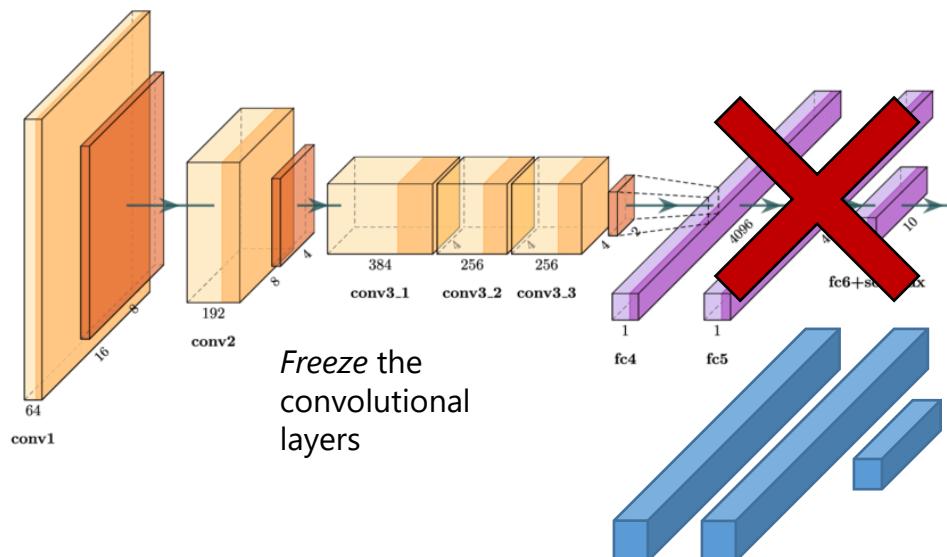
- `AlexNet`
- `VGG`
- `ResNet`
- `SqueezeNet`
- `DenseNet`
- `Inception v3`
- `GoogLeNet`
- `ShuffleNet v2`
- `MobileNetV2`
- `MobileNetV3`
- `ResNeXt`
- `Wide ResNet`
- `MNASNet`
- `EfficientNet`
- `RegNet`
- `VisionTransformer`
- `ConvNeXt`

1



Pretrained AlexNet
(pretrained on a lot of images)

2



Freeze the convolutional layers

3

Train the newly-added layers only.

Discard the last N layers.

Add new layers.

NB! There are several ways to do it depending on your dataset size. But not getting into details.

Transformer

CNN, RNN, Transformer

Transformer

```
CLASS torch.nn.TransformerEncoder(encoder_layer, num_layers,  
norm=None) [SOURCE]
```

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

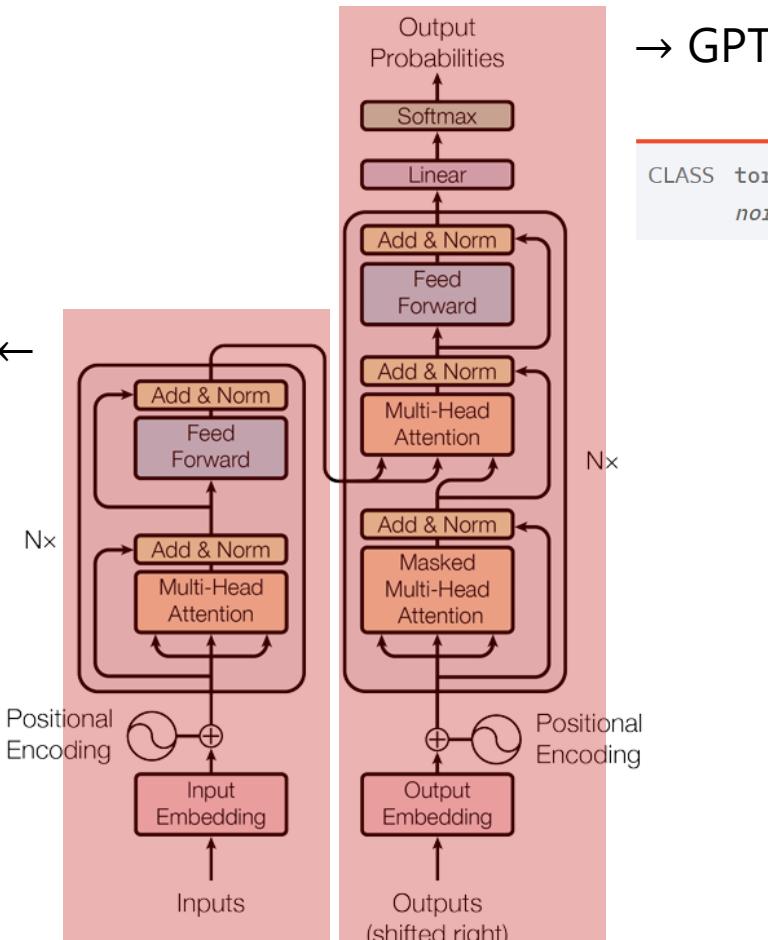
Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Łukasz Kaiser*
Google Brain
lukasz.kaiser@google.com

Illia Polosukhin* ‡
illia.polosukhin@gmail.com

BERT ←



→ GPT

```
CLASS torch.nn.TransformerDecoder(decoder_layer, num_layers,  
norm=None) [SOURCE]
```

CNN, RNN, Transformer

Transformer for Image

- Transformer for Image: Vision Transformer

PyTorch

Get Started Ecosystem Mobile Blog Tutorials Docs ▾

0.12 ▾

Search Docs

Package Reference

Transforming and augmenting images

Models and pre-trained weights

MODELS AND PRE-TRAINED WEIGHTS

The `torchvision.models` subpackage contains definitions of models for addressing different tasks including: image classification, pixelwise semantic segmentation, object detection, instance segmentation, person keypoint detection, video classification, and optical flow.

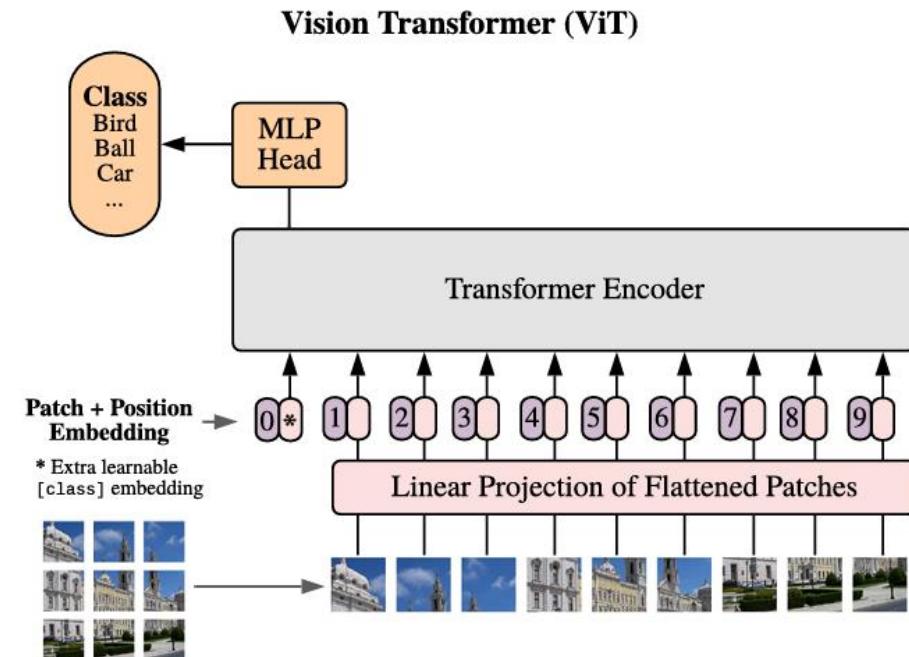
```
vit_b_16 = models.vit_b_16(pretrained=True)
vit_b_32 = models.vit_b_32(pretrained=True)
vit_l_16 = models.vit_l_16(pretrained=True)
vit_l_32 = models.vit_l_32(pretrained=True)
```

AN IMAGE IS WORTH 16x16 WORDS: TRANSFORMERS FOR IMAGE RECOGNITION AT SCALE

Alexey Dosovitskiy*,†, Lucas Beyer*, Alexander Kolesnikov*, Dirk Weissenborn*, Xiaohua Zhai*, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, Neil Houlsby*,†

*equal technical contribution, †equal advising

Google Research, Brain Team



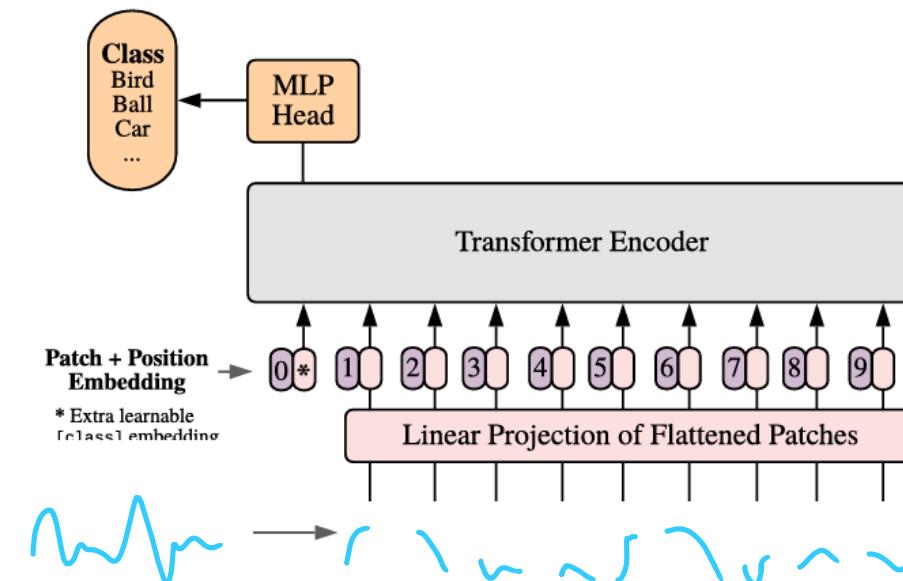
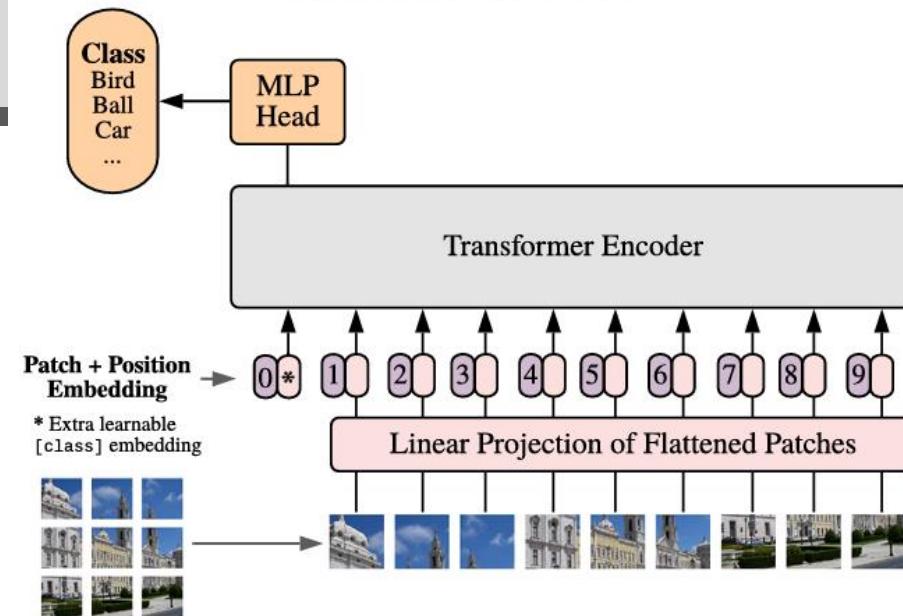
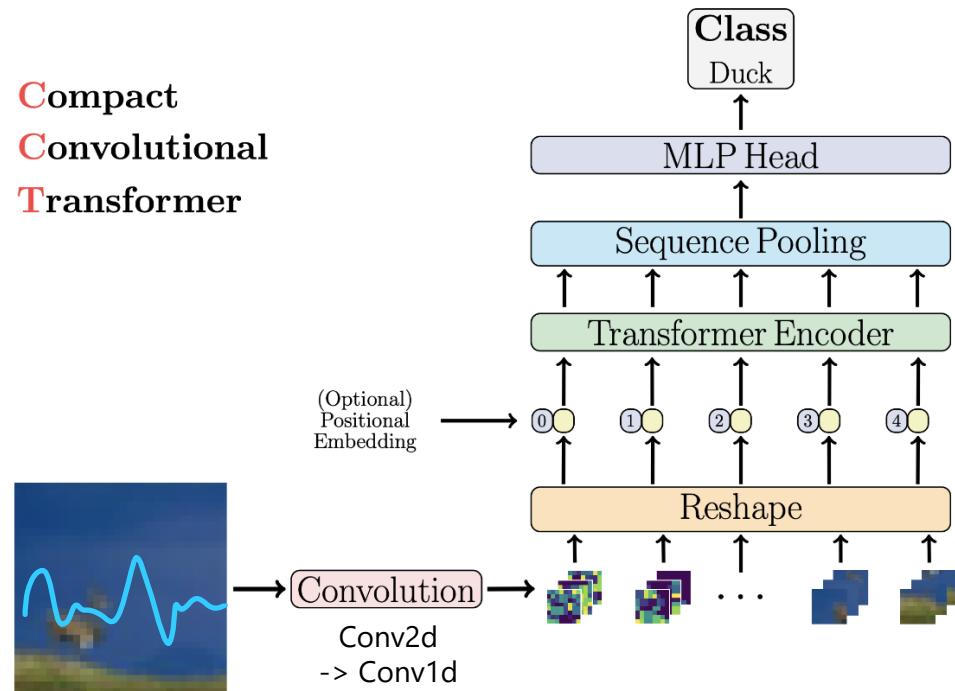
CNN, RNN, Transformer

Vision Transformer (ViT)

Transformer for Time Series

- Not much of standardized models.
- You can build one by yourself.

Compact
Convolutional
Transformer



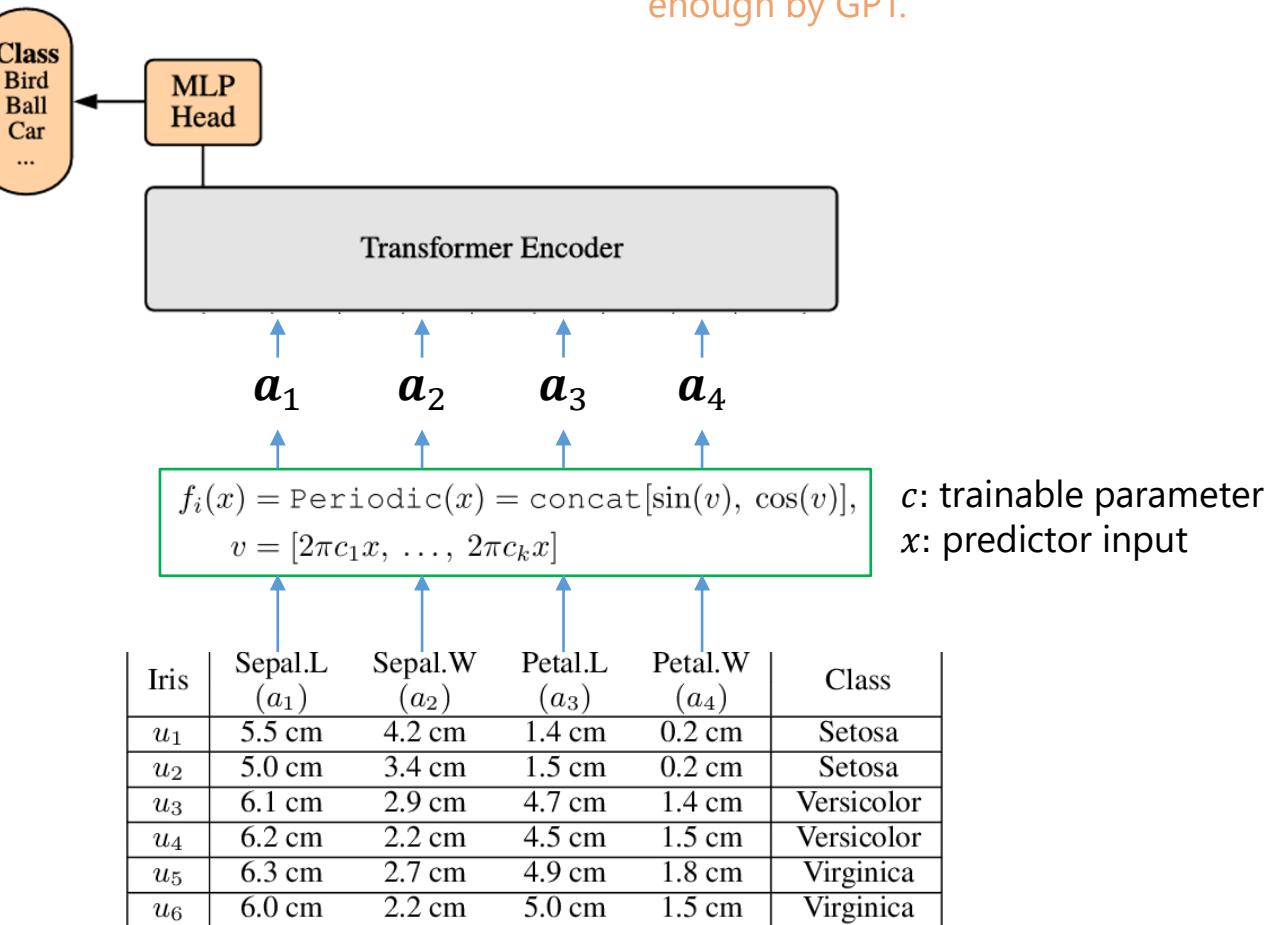
CNN, RNN, Transformer

Transformer for Tabular Data

- Relatively new research direction.
- Previous studies have shown that transformer models beat basic neural network models.
- Not much of standardized models.
- You can build one by yourself.

Try using the power of Transformer for your Tabular data 😊

The power of Transformer is proven enough by GPT.



Language Models: GPT, BERT

Language Models

Pretrained Language Models?

- e.g., Pretrained GPT

<https://www.crunchbase.com/organization/hugging-face> ::

Hugging Face - Crunchbase Company Profile & Funding

Hugging Face is an open-source & platform provider of machine learning technologies.

Hugging Face was launched in 2016 and is headquartered in New York.



Quick tour

```
from transformers import pipeline
classifier = pipeline("sentiment-analysis")
classifier("I'm so happy to do PhD at NTNU.")
[{'label': 'POSITIVE', 'score': 0.999579131603241}]
classifier("I paid so much tax this year.")
[{'label': 'NEGATIVE', 'score': 0.9869840145111084}]
```

The screenshot shows the Hugging Face website homepage. At the top, there's a search bar with the placeholder "Search models, datasets, users...". Below it, there are links for "Models", "Datasets", "Spaces", and "Docs". The main navigation bar has a "Transformers" section highlighted with a yellow smiley face emoji. Other sections include "GET STARTED" (with "Transformers" also highlighted), "Quick tour", "Installation", "TUTORIALS" (with "Pipelines for inference", "Load pretrained instances with an AutoClass", "Preprocess", "Fine-tune a pretrained model", "Distributed training with 🤗", and "Accelerate"), and "Changelog". To the right, there's a sidebar with the "Transformers" logo and the text: "State-of-the-art Machine Learning for PyTorch, TensorFlow and JAX. 🤗 Transformers provides APIs to easily download and train state-of-the-art pretrained models. Using pretrained models can reduce your compute costs, carbon footprint, and save you time from training a model from scratch. The models can be used across different modalities such as:

- 📄 Text: text classification, information extraction, question answering, summarization, translation, and text generation in over 100 languages.
- 🖼 Images: image classification, object detection, and segmentation.
- 🎧 Audio: speech recognition and audio classification.
- 🍃 Multimodal: table question answering, optical character recognition, information extraction from scanned documents, video classification, and visual question answering.

Language Models

Model for Different Tasks & Different Models

The screenshot shows the homepage of the Hugging Face Transformers documentation. At the top, there's a search bar with "Search documentation" and a "Ctrl+K" hotkey, followed by language selection dropdowns for "EN" and "V4.18.0". Below the header, there's a sidebar titled "TASKS" with a list of various NLP tasks: Text classification, Token classification, Question answering, Language modeling, Translation, Summarization, Multiple choice, Audio classification, Automatic speech recognition, and Image classification. The "Text classification" item is highlighted with a dark background. The main content area is mostly empty at the moment.

MODELS	DPT	PhoBERT	ViTMAE
ALBERT	ELECTRA	PLBart	VisualBERT
Auto Classes	Encoder Decoder Models	PoolFormer	Wav2Vec2
BART	FlauBERT	ProphetNet	Wav2Vec2Phoneme
BARThez	FNet	QDQBert	WavLM
BARTpho	FSMT	RAG	XGLM
BEiT	Funnel Transformer	REALM	XLM
BERT	GLPN	Reformer	XLM-ProphetNet
Bertweet	HerBERT	REMBERT	XLM-RoBERTa
BertGeneration	I-BERT	ResNet	XLM-RoBERTa-XL
BertJapanese	ImageGPT	RetriBERT	XLNet
BigBird	LayoutLM	RoBERTa	XLSR-Wav2Vec2
BigBirdPegasus	LayoutLMV2	RoFormer	XLS-R
Blenderbot	LayoutXLM	SegFormer	YOSO
Blenderbot Small	LED	SEW	
BORT	Longformer	SEW-D	
ByT5	LUKE	Speech Encoder Decoder Models	
CamemBERT	LXMERT	Speech2Text	
CANINE	MarianMT	Speech2Text2	
ConvNeXT	MaskFormer	Splinter	
CLIP	M2M100	SqueezeBERT	
ConvBERT	MBart and MBart-50	Swin Transformer	
CPM	MegatronBERT	T5	
CTRL	MegatronGPT2	T5v1.1	
Data2Vec	mLUKE	TAPAS	
DeBERTa	MobileBERT	Transformer XL	
DeBERTa-v2	MPNet	TrOCR	
Decision Transformer	MT5	UniSpeech	
DeiT	Nyströmlmformer	UniSpeech-SAT	
DETR	OpenAI GPT	VAN	
DialoGPT	OpenAI GPT2	ViLT	
DistilBERT	GPT-J	Vision Encoder Decoder Models	
	GPT Neo	Vision Text Dual Encoder	
DiT	Hubert	Vision Transformer (ViT)	
DPR	Perceiver	ViTMAE	
DPT	Pegasus		

Recommend Study Materials

Recommend Study Materials

- Deep Neural Networks with PyTorch in COURSERA
- Examples from PyTorch official documentation
- StackOverFlow
- etc.

Browse > Data Science > Machine Learning

This course is part of the IBM AI Engineering Professional Certificate

Deep Neural Networks with PyTorch

★★★★★ 4.4 1,086 ratings | ⚡ 88%



Joseph Santarcangelo

Offered By 

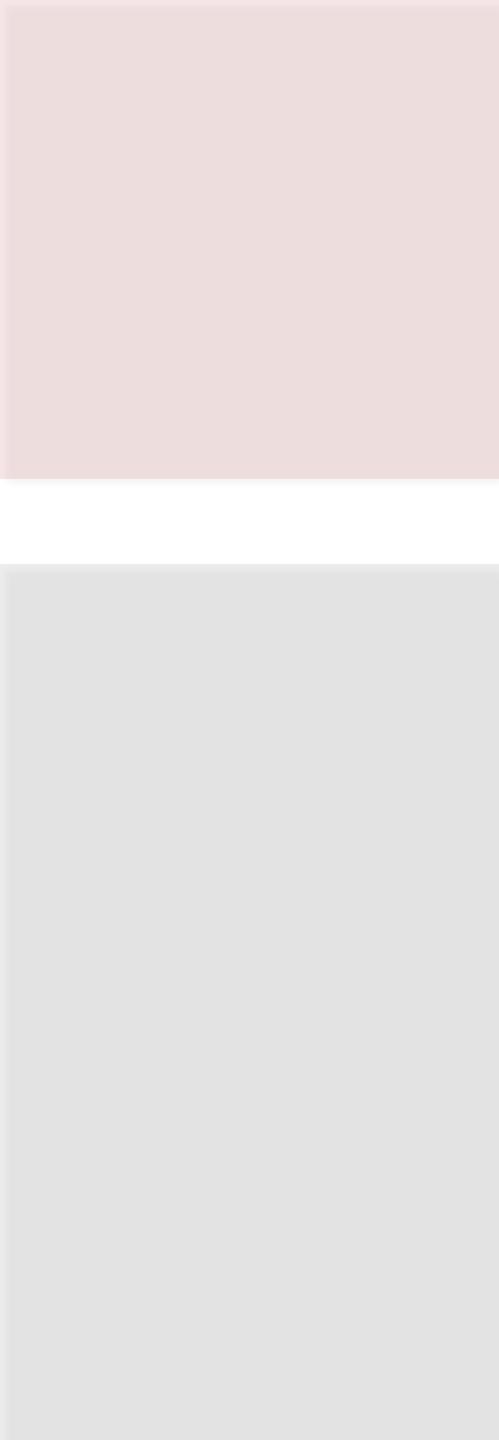
 PyTorch

Get Started Ecosystem Mobile Blog Tutorials Docs

1.11.0+cu102

Tutorials > Welcome to PyTorch Tutorials

WELCOME TO PYTORCH TUTORIALS



Highly Recommended Tools and Libraries

(that can make your life easier.)

Highly Recommended Tools and Libraries

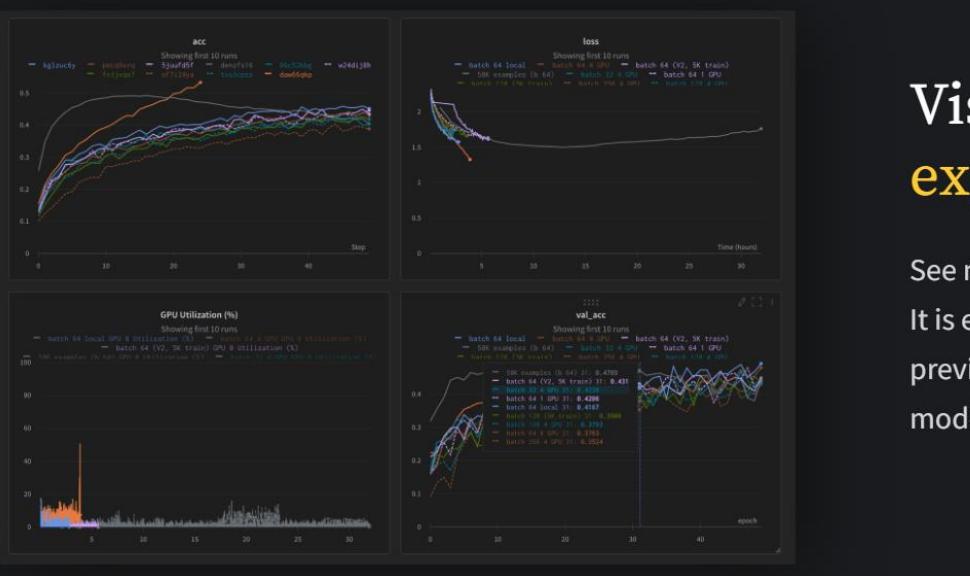
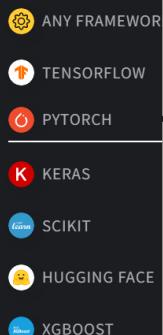
- **Weights and Biases (wandb)**: cloud-based result tracking tool
- **PyTorch lightning**: simplified version of PyTorch

Highly Recommended

Weights and Biases (wandb)

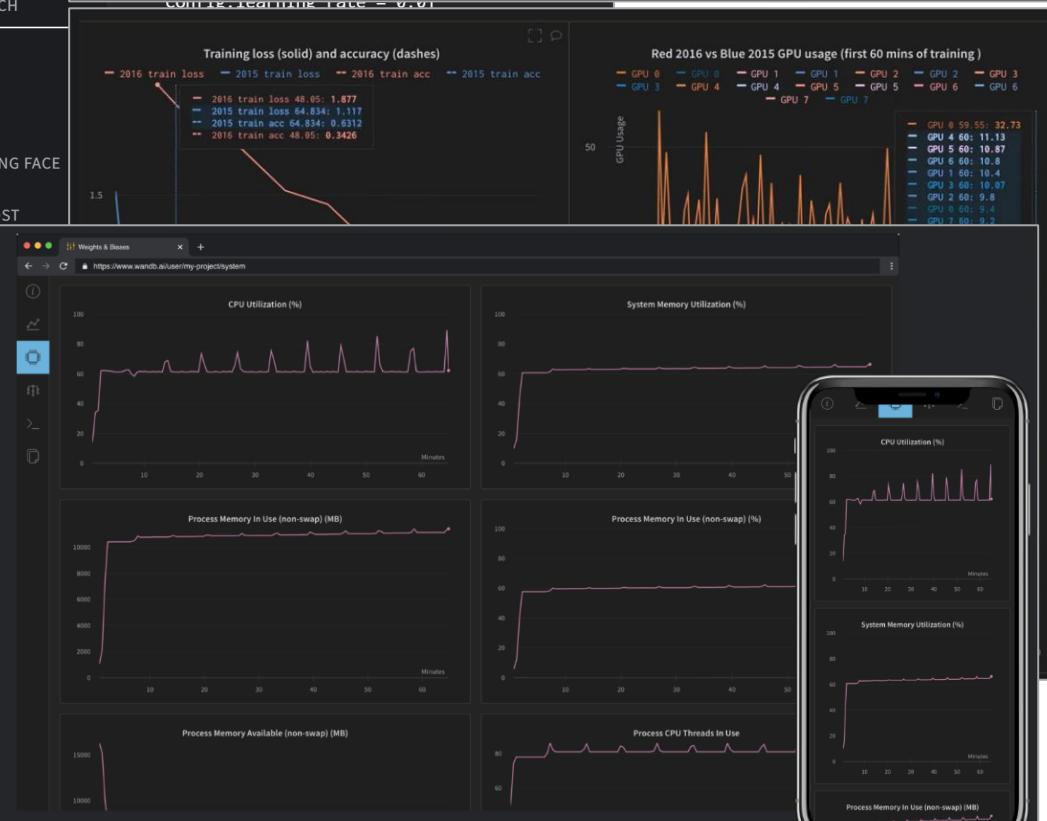
Track, compare, and visualize with 5 lines of code

Add a few lines to your script to start logging results.
Our lightweight integration works with any Python script.



Accessible anywhere

Check the latest training model and results on desktop and mobile. Use collaborative hosted projects to coordinate across your team.



Visualize and compare every experiment

See model metrics stream live into interactive graphs and tables. It is easy to see how your latest model is performing compared to previous experiments, no matter where you are training your models.

Monitor your CPU and GPU usage

Visualize live metrics like GPU utilization to identify training bottlenecks and avoid wasting expensive resources.

UCR-ALL-kNN Workspace – Weig +

wandb.ai/daesoolee/UCR-ALL-kNN?workspace=user-daesoolee

MyBlog G Cloud NTNU Norway HPC Arxiv Project

daesoolee > Projects > UCR-ALL-kNN

Runs (1878)

Search panels

Charts 24

simclr_loss.test simclr_loss.train simclr_loss.validate kNN_acc var_loss.train

feature_decorr_metrics sim_loss.train test_loss sim_loss.test loss.test

Media 8

partial(abs(C))-ep_100 partial(abs(C))-ep_40

1-10 of 24 1-2 of 8

53

NonInvasiveFetalECGThorax1-barlow_tv 0.8254
 NonInvasiveFetalECGThorax1-byol 0.7094
 NonInvasiveFetalECGThorax1-simclr 0.7547
 NonInvasiveFetalECGThorax1-simsiam 0.8601
 NonInvasiveFetalECGThorax1-vibcreg 0.5064
 NonInvasiveFetalECGThorax1-vibcreg_si 0.5079
 NonInvasiveFetalECGThorax1-vibcreg_si 0.8244
 NonInvasiveFetalECGThorax1-vibcreg_si 0.8427
 NonInvasiveFetalECGThorax1-vibcreg_si 0.8529
 NonInvasiveFetalECGThorax1-vibcreg_si 0.8051
 NonInvasiveFetalECGThorax1-vibcreg_si 0.8193
 NonInvasiveFetalECGThorax1-vibcreg_si 0.8529
 NonInvasiveFetalECGThorax1-vibcreg_si 0.7374
 NonInvasiveFetalECGThorax1-vibcreg_si -
 NonInvasiveFetalECGThorax1-vibcreg_si 0.7552
 NonInvasiveFetalECGThorax1-vibcreg_si 0.7506
 NonInvasiveFetalECGThorax1-vibcreg_si 0.6265
 NonInvasiveFetalECGThorax1-vibcreg_si 0.5257
 NonInvasiveFetalECGThorax1-vibcreg_si 0.6061
 NonInvasiveFetalECGThorax1-vibcreg_si 0.6967

Weights & Biases Products Resources Docs Pricing Enterprise

Articles Projects ML News Events Podcast Forum

W&B Fully Connected: Articles > The Science of Debugging with W&B Reports

Highly Recommended To

Weights and Biases (wandb)

The Science of Debugging with W&B Reports

Sarah Jane



Last Updated: Feb 7, 2022

Overview

Debugging Inside a Research Organization

At Latent Space, we're pushing the frontier of generative modeling with a focused group of researchers and engineer. But like anytime you're leading novel research, chances are you're going to run into some novel challenges too. This report will walk you through how we use W&B reports to help us debug models inside a research organization. Let's get started.

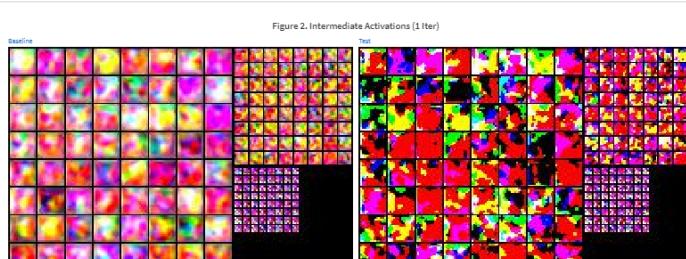
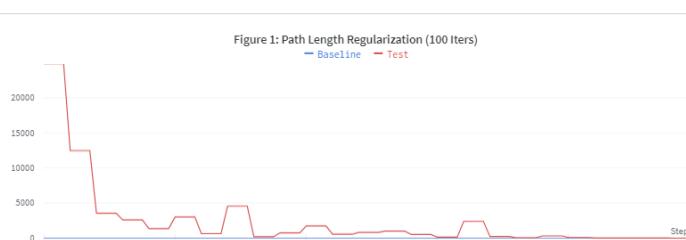
Challenge

Bug: Evident at the start of training (100 iterations)

At LatentSpace, we train an order of magnitude more models than is possible to humanly debug. In one instance, we modified a module in test which should not have affected training quality, in fact, we know prior to kicking off the test (red) and baseline (blue) should be equivalent. However, we quickly noticed that the training dynamics changed – looks like a bug!

This was evident when monitoring one of the regularization terms. As seen in figure 1, the problem is evident since the test run is a magnitude order larger than the baseline.

In figure 2, the test run looks more saturated when compared to the baseline. This is quite suspicious.



Highly Recorded

3D Visualizations

test/examples

3D Object Point Clouds Molecules

Log files in the formats 'obj', 'gltf', 'glb', 'babylon', 'stl', 'pts.json', and we will render them in the UI when your run finishes.

```
1 wandb.log({"generated_samples":  
2     [wandb.Object3D(open("sample.obj")),  
3      wandb.Object3D(open("sample.gltf")),  
4      wandb.Object3D(open("sample.glb"))]})
```

ries

Other Media

Weights & Biases also supports logging of a variety of other media types.

Audio Video Text HTML

```
1 wandb.log(  
2 {"whale_songs": wandb.Audio(np_array, caption="OooOoo", sample_rate=32)})
```

The maximum number of audio clips that can be logged per step is 100.

Segmentation Masks Bounding Boxes

Log semantic segmentation masks and interact with them (altering opacity, viewing changes over time, and more) via the W&B UI.

Index 0

ground truth

road sidewalk building wall fence pole traffic light tra

Prediction improvement over time



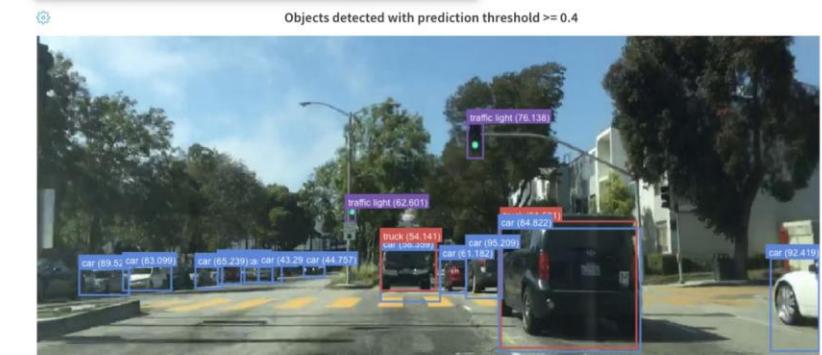
Segmentation Masks Bounding Boxes

Log bounding boxes with images, and use filters and toggles to dynamically visualize different sets of boxes in the UI.

score \geq 40.4

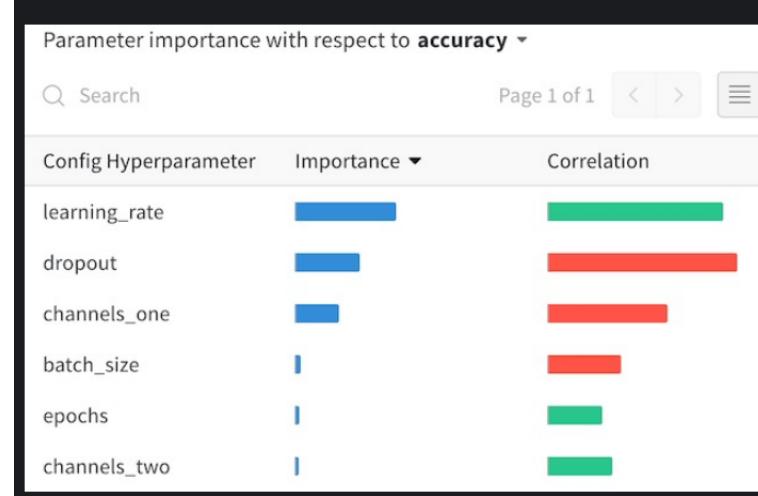
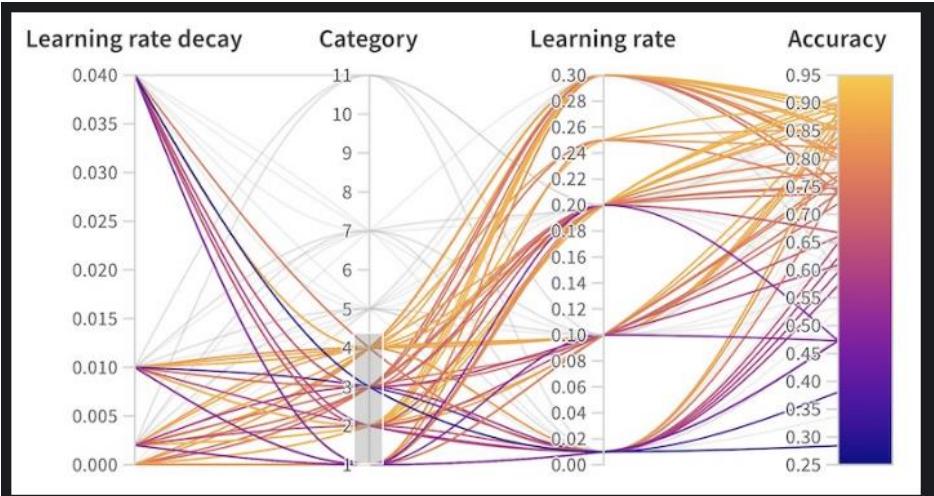
predictions

car truck person traffic light stop sign bus bicycle mot



Highly Recommended Tools and Libraries

Weights and Biases (wandb)

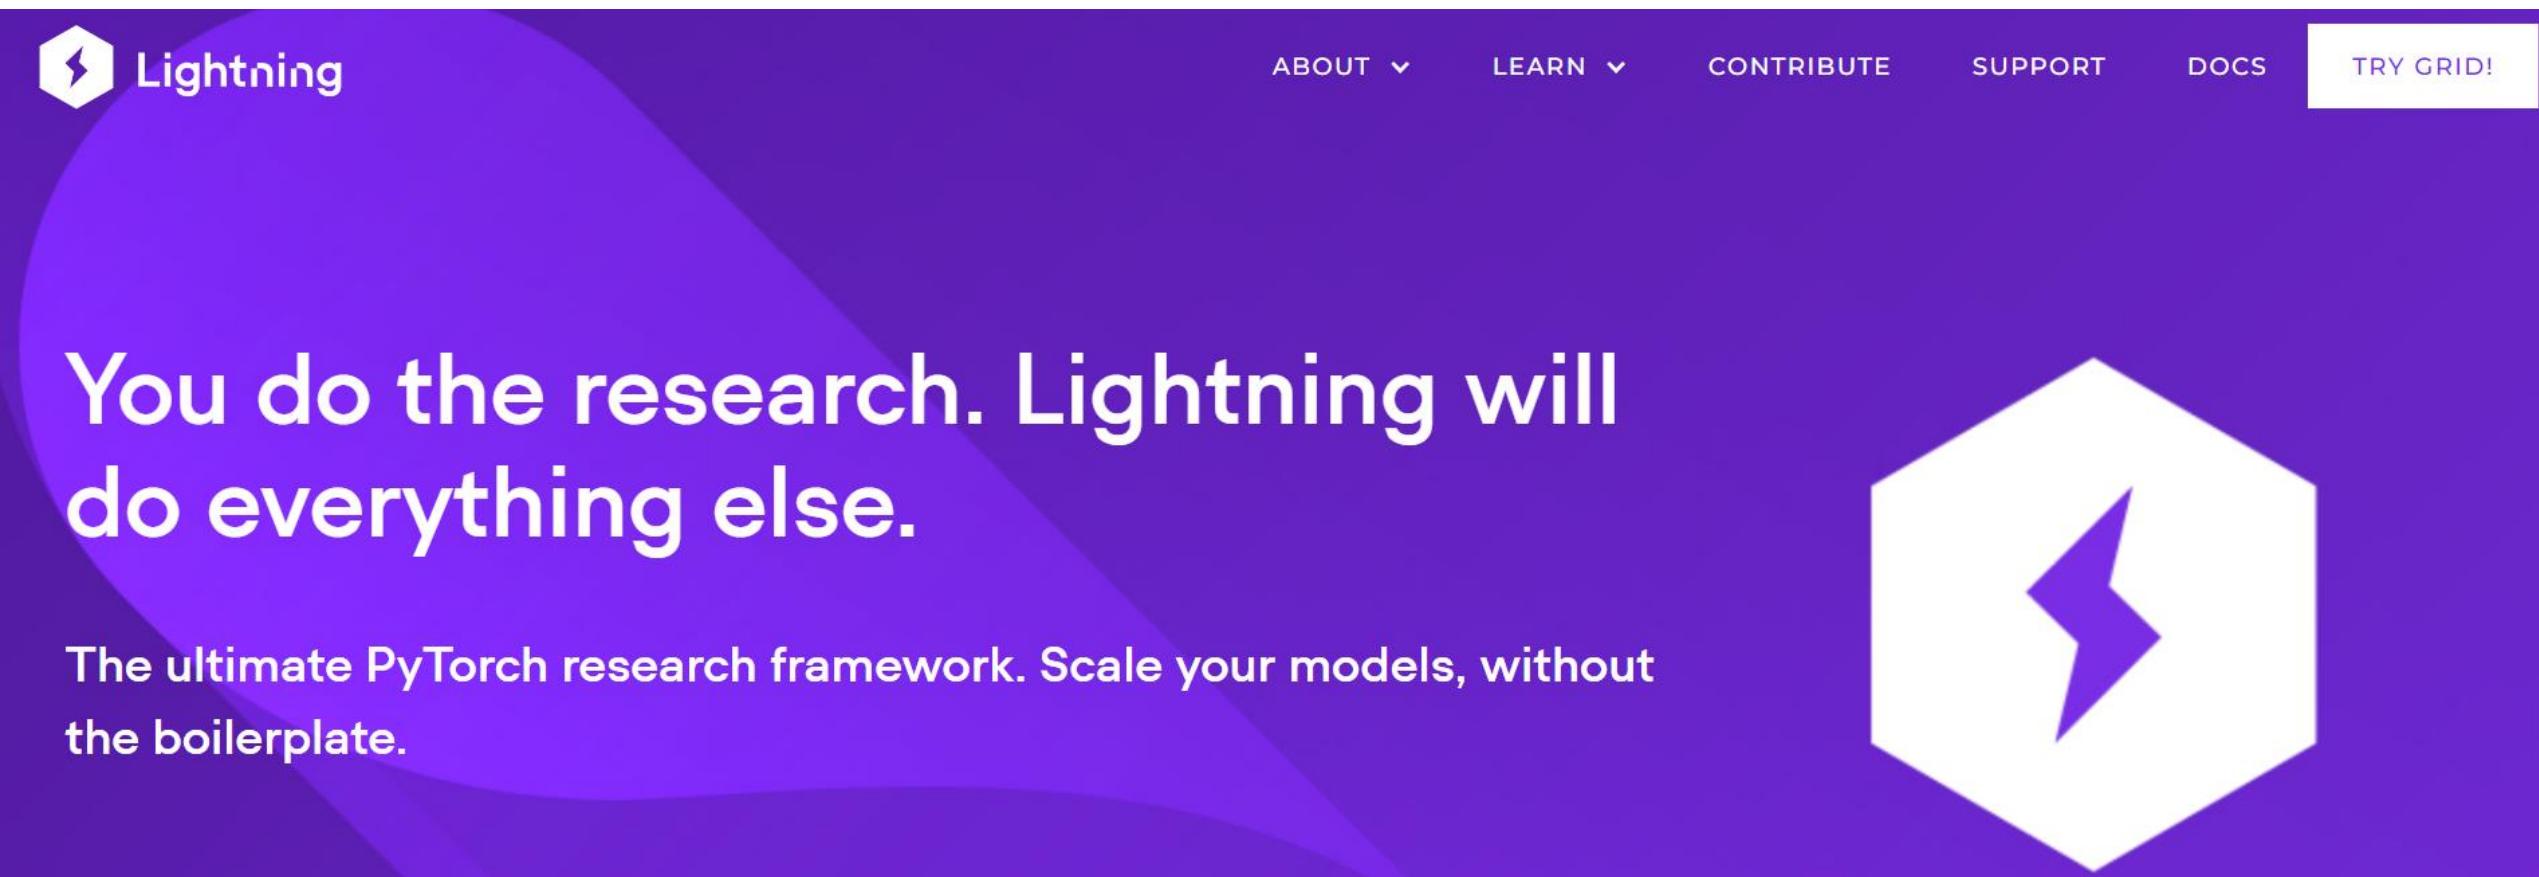


Parameter importance

Visualize which hyperparameters affect the metrics you care about. W&B comes with default visualizations that make it easy to get started without writing custom code to compare experiments.

Highly Recommended Tools and Libraries

PyTorch lightning: simplified version of PyTorch

The image shows the homepage of PyTorch Lightning. The background is a dark purple color with a large white hexagonal logo in the center. The logo contains a stylized purple lightning bolt. At the top left, there is a smaller version of the logo next to the word "Lightning". At the top right, there is a navigation bar with links: "ABOUT", "LEARN", "CONTRIBUTE", "SUPPORT", "DOCS", and a button labeled "TRY GRID!". Below the logo, the main slogan reads "You do the research. Lightning will do everything else." in a large, bold, white font. Below this, a subtitle in a smaller white font reads "The ultimate PyTorch research framework. Scale your models, without the boilerplate.".

Highly Recommended

PyTorch

PyTorch lightning:

```
# model
class Net(nn.Module):
    def __init__(self):
        self.layer_1 = torch.nn.Linear(28 * 28, 128)
        self.layer_2 = torch.nn.Linear(128, 10)

    def forward(self, x):
        x = x.view(x.size(0), -1)
        x = self.layer_1(x)
        x = F.relu(x)
        x = self.layer_2(x)
        return x

    # train loader
mnist_train = MNIST(os.getcwd(), train=True, download=True,
                     transform=transforms.ToTensor())
mnist_train = DataLoader(mnist_train, batch_size=64)

net = Net()

# optimizer + scheduler
optimizer = torch.optim.Adam(net.parameters(), lr=1e-3)
scheduler = StepLR(optimizer, step_size=1)

# train
for epoch in range(1, 100):
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        output = model(data)
        loss = F.nll_loss(output, target)

        loss.backward()
        optimizer.step()
        if batch_idx % args.log_interval == 0:
            print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
                epoch, batch_idx * len(data), len(train_loader.dataset),
                100. * batch_idx / len(train_loader), loss.item()))
```

PyTorch Lightning

```
# model
class Net(LightningModule):
    def __init__(self):
        self.layer_1 = torch.nn.Linear(28 * 28, 128)
        self.layer_2 = torch.nn.Linear(128, 10)

    def forward(self, x):
        x = x.view(x.size(0), -1)
        x = self.layer_1(x)
        x = F.relu(x)
        x = self.layer_2(x)
        return x

    def train_dataloader(self):
        mnist_train = MNIST(os.getcwd(), train=True, download=True,
                           transform=transforms.ToTensor())
        return DataLoader(mnist_train, batch_size=64)

    def configure_optimizers(self):
        optimizer = torch.optim.Adam(self.parameters(), lr=1e-3)
        scheduler = StepLR(optimizer, step_size=1)
        return optimizer, scheduler

    def training_step(self, batch, batch_idx):
        data, target = batch
        output = self.forward(data)
        loss = F.nll_loss(output, target)
        return loss
```

NB! Training code becomes significantly simpler.

Thank you!



Norwegian University of
Science and Technology

Why PyTorch?

When to use Tensorflow?

